

[INFO] 1 Control Plane Security Configuration

[INFO] 1.1 Control Plane Node Configuration Files

[FAIL] 1.1.1 Ensure that the API server pod specification file permissions are set to 600 or more restrictive (Automated)

[FAIL] 1.1.2 Ensure that the API server pod specification file ownership is set to root:root (Automated)

[FAIL] 1.1.3 Ensure that the controller manager pod specification file permissions are set to 600 or more restrictive (Automated)

[FAIL] 1.1.4 Ensure that the controller manager pod specification file ownership is set to root:root (Automated)

[FAIL] 1.1.5 Ensure that the scheduler pod specification file permissions are set to 600 or more restrictive (Automated)

[FAIL] 1.1.6 Ensure that the scheduler pod specification file ownership is set to root:root (Automated)

[FAIL] 1.1.7 Ensure that the etcd pod specification file permissions are set to 600 or more restrictive (Automated)

[FAIL] 1.1.8 Ensure that the etcd pod specification file ownership is set to root:root (Automated)

[WARN] 1.1.9 Ensure that the Container Network Interface file permissions are set to 600 or more restrictive (Manual)

[WARN] 1.1.10 Ensure that the Container Network Interface file ownership is set to root:root (Manual)

[FAIL] 1.1.11 Ensure that the etcd data directory permissions are set to 700 or more restrictive (Automated)

[FAIL] 1.1.12 Ensure that the etcd data directory ownership is set to etcd:etcd (Automated)

[FAIL] 1.1.13 Ensure that the admin.conf file permissions are set to 600 or more restrictive (Automated)

[FAIL] 1.1.14 Ensure that the admin.conf file ownership is set to root:root (Automated)

[FAIL] 1.1.15 Ensure that the scheduler.conf file permissions are set to 600 or more restrictive (Automated)

[FAIL] 1.1.16 Ensure that the scheduler.conf file ownership is set to root:root (Automated)

[FAIL] 1.1.17 Ensure that the controller-manager.conf file permissions are set to 600 or more restrictive (Automated)

[FAIL] 1.1.18 Ensure that the controller-manager.conf file ownership is set to root:root (Automated)

[FAIL] 1.1.19 Ensure that the Kubernetes PKI directory and file ownership is set to root:root (Automated)

[WARN] 1.1.20 Ensure that the Kubernetes PKI certificate file permissions are set to 600 or more restrictive (Manual)

[WARN] 1.1.21 Ensure that the Kubernetes PKI key file permissions are set to 600 (Manual)

[INFO] 1.2 API Server

[PASS] 1.2.1 Ensure that the --anonymous-auth argument is set to false (Manual)

[PASS] 1.2.2 Ensure that the --token-auth-file parameter is not set (Automated)

[WARN] 1.2.3 Ensure that the --DenyServiceExternalIPs is set (Manual)

[PASS] 1.2.4 Ensure that the --kubelet-client-certificate and --kubelet-client-key arguments are set as appropriate (Automated)

[PASS] 1.2.5 Ensure that the --kubelet-certificate-authority argument is set as appropriate (Automated)

[PASS] 1.2.6 Ensure that the --authorization-mode argument is not set to AlwaysAllow (Automated)

[PASS] 1.2.7 Ensure that the --authorization-mode argument includes Node (Automated)

[PASS] 1.2.8 Ensure that the --authorization-mode argument includes RBAC (Automated)

[WARN] 1.2.9 Ensure that the admission control plugin EventRateLimit is set (Manual)

[PASS] 1.2.10 Ensure that the admission control plugin AlwaysAdmit is not set (Automated)

[WARN] 1.2.11 Ensure that the admission control plugin AlwaysPullImages is set (Manual)

[WARN] 1.2.12 Ensure that the admission control plugin SecurityContextDeny is set if PodSecurityPolicy is not used (Manual)

[PASS] 1.2.13 Ensure that the admission control plugin ServiceAccount is set (Automated)

[PASS] 1.2.14 Ensure that the admission control plugin NamespaceLifecycle is set (Automated)

[PASS] 1.2.15 Ensure that the admission control plugin NodeRestriction is set (Automated)

[PASS] 1.2.16 Ensure that the --secure-port argument is not set to 0 - NoteThis recommendation is obsolete and will be deleted per the consensus process (Manual)

[PASS] 1.2.17 Ensure that the --profiling argument is set to false (Automated)

[PASS] 1.2.18 Ensure that the --audit-log-path argument is set (Automated)

[PASS] 1.2.19 Ensure that the --audit-log-maxage argument is set to 30 or as appropriate (Automated)

[PASS] 1.2.20 Ensure that the --audit-log-maxbackup argument is set to 10 or as appropriate (Automated)

[FAIL] 1.2.21 Ensure that the --audit-log-maxsize argument is set to 100 or as appropriate (Automated)

[WARN] 1.2.22 Ensure that the --request-timeout argument is set as appropriate (Manual)

[PASS] 1.2.23 Ensure that the --service-account-lookup argument is set to true (Automated)

[PASS] 1.2.24 Ensure that the --service-account-key-file argument is set as appropriate (Automated)

[PASS] 1.2.25 Ensure that the --etcd-certfile and --etcd-keyfile arguments are set as appropriate (Automated)

[PASS] 1.2.26 Ensure that the --tls-cert-file and --tls-private-key-file arguments are set as appropriate (Automated)

[PASS] 1.2.27 Ensure that the --client-ca-file argument is set as appropriate (Automated)

[PASS] 1.2.28 Ensure that the --etcd-cafile argument is set as appropriate (Automated)

[WARN] 1.2.29 Ensure that the --encryption-provider-config argument is set as appropriate (Manual)

[WARN] 1.2.30 Ensure that encryption providers are appropriately configured (Manual)

[PASS] 1.2.31 Ensure that the API Server only makes use of Strong Cryptographic Ciphers (Manual)

[INFO] 1.3 Controller Manager

[WARN] 1.3.1 Ensure that the --terminated-pod-gc-threshold argument is set as appropriate (Manual)

[PASS] 1.3.2 Ensure that the --profiling argument is set to false (Automated)

[PASS] 1.3.3 Ensure that the --use-service-account-credentials argument is set to true (Automated)

[PASS] 1.3.4 Ensure that the --service-account-private-key-file argument is set as appropriate (Automated)

[PASS] 1.3.5 Ensure that the --root-ca-file argument is set as appropriate (Automated)

[PASS] 1.3.6 Ensure that the RotateKubeletServerCertificate argument is set to true (Automated)

[PASS] 1.3.7 Ensure that the --bind-address argument is set to 127.0.0.1 (Automated)

[INFO] 1.4 Scheduler

[PASS] 1.4.1 Ensure that the --profiling argument is set to false (Automated)

[PASS] 1.4.2 Ensure that the --bind-address argument is set to 127.0.0.1 (Automated)

== Remediations master ==

1.1.1 Run the below command (based on the file location on your system) on the control plane node.

For example, `chmod 600 /etc/pf9/kube.d/pod-manifests/master.yaml`

1.1.2 Run the below command (based on the file location on your system) on the control plane node.

For example, `chown root:root /etc/pf9/kube.d/pod-manifests/master.yaml`

1.1.3 Run the below command (based on the file location on your system) on the control plane node.

For example, `chmod 600 /etc/pf9/kube.d/pod-manifests/master.yaml`

1.1.4 Run the below command (based on the file location on your system) on the control plane node.

For example, `chown root:root /etc/pf9/kube.d/pod-manifests/master.yaml`

1.1.5 Run the below command (based on the file location on your system) on the control plane node.

For example, `chmod 600 /etc/pf9/kube.d/pod-manifests/master.yaml`

1.1.6 Run the below command (based on the file location on your system) on the control plane node.

For example, `chown root:root /etc/pf9/kube.d/pod-manifests/master.yaml`

1.1.7 Run the below command (based on the file location on your system) on the control plane node.

For example,

`chmod 600 /etc/kubernetes/manifests/etcd.yaml`

1.1.8 Run the below command (based on the file location on your system) on the control plane node.

For example,

`chown root:root /etc/kubernetes/manifests/etcd.yaml`

1.1.9 Run the below command (based on the file location on your system) on the control plane node.

For example, `chmod 600 <path/to/cni/files>`

1.1.10 Run the below command (based on the file location on your system) on the control plane node.

For example,

`chown root:root <path/to/cni/files>`

1.1.11 On the etcd server node, get the etcd data directory, passed as an argument `--data-dir`, from the command `'ps -ef | grep etcd'`.

Run the below command (based on the etcd data directory found above). For example,

`chmod 700 /var/lib/etcd`

1.1.12 On the etcd server node, get the etcd data directory, passed as an argument `--data-dir`, from the command `'ps -ef | grep etcd'`.

Run the below command (based on the etcd data directory found above).

For example, `chown etcd:etcd /var/lib/etcd`

1.1.13 Run the below command (based on the file location on your system) on the control plane node.

For example, `chmod 600 /etc/kubernetes/admin.conf`

1.1.14 Run the below command (based on the file location on your system) on the control plane node.

For example, `chown root:root /etc/kubernetes/admin.conf`

1.1.15 Run the below command (based on the file location on your system) on the control plane node.

For example,

`chmod 600 /etc/kubernetes/scheduler.conf`

1.1.16 Run the below command (based on the file location on your system) on the control plane node.

For example,

`chown root:root /etc/kubernetes/scheduler.conf`

1.1.17 Run the below command (based on the file location on your system) on the control plane node.

For example,

`chmod 600 /etc/kubernetes/controller-manager.conf`

1.1.18 Run the below command (based on the file location on your system) on the control plane node.

For example,

`chown root:root /etc/kubernetes/controller-manager.conf`

1.1.19 Run the below command (based on the file location on your system) on the control plane node.

For example,

`chown -R root:root /etc/kubernetes/pki/`

1.1.20 audit test did not run: failed to run: "find /etc/kubernetes/pki/ -name '*.crt' | xargs stat -c permissions=%a", output: "find: '/etc/kubernetes/pki/': No such file or directory\nstat: missing operand\nTry 'stat --help' for more information.\n", error: exit status 123

1.1.21 audit test did not run: failed to run: "find /etc/kubernetes/pki/ -name '*.key' | xargs stat -c permissions=%a", output: "find: '/etc/kubernetes/pki/': No such file or directory\nstat: missing operand\nTry 'stat --help' for more information.\n", error: exit status 123

1.2.3 Edit the API server pod specification file `/etc/pf9/kube.d/pod-manifests/master.yaml` on the control plane node and remove the `'DenyServiceExternalIPs'` from enabled admission plugins.

1.2.9 Follow the Kubernetes documentation and set the desired limits in a configuration file. Then, edit the API server pod specification file `/etc/pf9/kube.d/pod-manifests/master.yaml` and set the below parameters.

`--enable-admission-plugins=...,EventRateLimit,...`

`--admission-control-config-file=<path/to/configuration/file>`

1.2.11 Edit the API server pod specification file `/etc/pf9/kube.d/pod-manifests/master.yaml` on the control plane node and set the `--enable-admission-plugins` parameter to include `AlwaysPullImages`.

`--enable-admission-plugins=...,AlwaysPullImages,...`

1.2.12 Edit the API server pod specification file `/etc/pf9/kube.d/pod-manifests/master.yaml` on the control plane node and set the `--enable-admission-plugins` parameter to include `SecurityContextDeny`, unless `PodSecurityPolicy` is already in place.

`--enable-admission-plugins=...,SecurityContextDeny,...`

1.2.21 Edit the API server pod specification file `/etc/pf9/kube.d/pod-manifests/master.yaml` on the control plane node and set the `--audit-log-maxsize` parameter to an appropriate size in MB. For example, to set it as 100 MB, `--audit-log-maxsize=100`

1.2.22 Edit the API server pod specification file `/etc/pf9/kube.d/pod-manifests/master.yaml` and set the below parameter as appropriate and if needed.

For example, `--request-timeout=300s`

1.2.29 Follow the Kubernetes documentation and configure a `EncryptionConfig` file.

Then, edit the API server pod specification file `/etc/pf9/kube.d/pod-manifests/master.yaml` on the control plane node and set the `--encryption-provider-config` parameter to the path of that file.

For example, `--encryption-provider-config=</path/to/EncryptionConfig/File>`

1.2.30 Follow the Kubernetes documentation and configure a `EncryptionConfig` file.

In this file, choose aesCBC, kms or secretbox as the encryption provider.

1.3.1 Edit the Controller Manager pod specification file /etc/pf9/kube.d/pod-manifests/master.yaml on the control plane node and set the --terminated-pod-gc-threshold to an appropriate threshold, for example, --terminated-pod-gc-threshold=10

== Summary master ==

31 checks PASS

18 checks FAIL

12 checks WARN

0 checks INFO

[INFO] 2 Etcd Node Configuration

[INFO] 2 Etcd Node Configuration

[PASS] 2.1 Ensure that the --cert-file and --key-file arguments are set as appropriate (Automated)

[PASS] 2.2 Ensure that the --client-cert-auth argument is set to true (Automated)

[PASS] 2.3 Ensure that the --auto-tls argument is not set to true (Automated)

[PASS] 2.4 Ensure that the --peer-cert-file and --peer-key-file arguments are set as appropriate (Automated)

[PASS] 2.5 Ensure that the --peer-client-cert-auth argument is set to true (Automated)

[PASS] 2.6 Ensure that the --peer-auto-tls argument is not set to true (Automated)

[PASS] 2.7 Ensure that a unique Certificate Authority is used for etcd (Manual)

== Summary etcd ==

7 checks PASS

0 checks FAIL

0 checks WARN

0 checks INFO

[INFO] 3 Control Plane Configuration

[INFO] 3.1 Authentication and Authorization

[WARN] 3.1.1 Client certificate authentication should not be used for users (Manual)

[WARN] 3.1.2 Service account token authentication should not be used for users (Manual)

[WARN] 3.1.3 Bootstrap token authentication should not be used for users (Manual)

[INFO] 3.2 Logging

[WARN] 3.2.1 Ensure that a minimal audit policy is created (Manual)

[WARN] 3.2.2 Ensure that the audit policy covers key security concerns (Manual)

== Remediations controlplane ==

3.1.1 Alternative mechanisms provided by Kubernetes such as the use of OIDC should be implemented in place of client certificates.

3.1.2 Alternative mechanisms provided by Kubernetes such as the use of OIDC should be implemented

in place of service account tokens.

3.1.3 Alternative mechanisms provided by Kubernetes such as the use of OIDC should be implemented

in place of bootstrap tokens.

3.2.1 Create an audit policy file for your cluster.

3.2.2 Review the audit policy provided for the cluster and ensure that it covers at least the following areas,

- Access to Secrets managed by the cluster. Care should be taken to only log Metadata for requests to Secrets, ConfigMaps, and TokenReviews, in

order to avoid risk of logging sensitive data.

- Modification of Pod and Deployment objects.

- Use of `pods/exec`, `pods/portforward`, `pods/proxy` and `services/proxy`.

For most requests, minimally logging at the Metadata level is recommended (the most basic level of logging).

== Summary controlplane ==

0 checks PASS

0 checks FAIL

5 checks WARN

0 checks INFO

[INFO] 4 Worker Node Security Configuration

[INFO] 4.1 Worker Node Configuration Files

[FAIL] 4.1.1 Ensure that the kubelet service file permissions are set to 600 or more restrictive (Automated)

[FAIL] 4.1.2 Ensure that the kubelet service file ownership is set to root:root (Automated)

[PASS] 4.1.3 If proxy kubeconfig file exists ensure permissions are set to 600 or more restrictive (Manual)

[PASS] 4.1.4 If proxy kubeconfig file exists ensure ownership is set to root:root (Manual)

[PASS] 4.1.5 Ensure that the --kubeconfig kubelet.conf file permissions are set to 600 or more restrictive (Automated)

[FAIL] 4.1.6 Ensure that the --kubeconfig kubelet.conf file ownership is set to root:root (Automated)

[WARN] 4.1.7 Ensure that the certificate authorities file permissions are set to 600 or more restrictive (Manual)

[WARN] 4.1.8 Ensure that the client certificate authorities file ownership is set to root:root (Manual)

[WARN] 4.1.9 If the kubelet config.yaml configuration file is being used validate permissions set to 600 or more restrictive (Manual)

[WARN] 4.1.10 If the kubelet config.yaml configuration file is being used validate file ownership is set to root:root (Manual)

[INFO] 4.2 Kubelet

[PASS] 4.2.1 Ensure that the --anonymous-auth argument is set to false (Automated)

[FAIL] 4.2.2 Ensure that the --authorization-mode argument is not set to AlwaysAllow (Automated)

[PASS] 4.2.3 Ensure that the --client-ca-file argument is set as appropriate (Automated)

[PASS] 4.2.4 Verify that the --read-only-port argument is set to 0 (Manual)

[PASS] 4.2.5 Ensure that the --streaming-connection-idle-timeout argument is not set to 0 (Manual)

[PASS] 4.2.6 Ensure that the --make-iptables-util-chains argument is set to true (Automated)

[WARN] 4.2.7 Ensure that the --hostname-override argument is not set (Manual)

[PASS] 4.2.8 Ensure that the eventRecordQPS argument is set to a level which ensures appropriate event capture (Manual)

[PASS] 4.2.9 Ensure that the --tls-cert-file and --tls-private-key-file arguments are set as appropriate (Manual)

[PASS] 4.2.10 Ensure that the --rotate-certificates argument is not set to false (Automated)

[PASS] 4.2.11 Verify that the RotateKubeletServerCertificate argument is set to true (Manual)

[PASS] 4.2.12 Ensure that the Kubelet only makes use of Strong Cryptographic Ciphers (Manual)

[WARN] 4.2.13 Ensure that a limit is set on pod PIDs (Manual)

== Remediations node ==

4.1.1 Run the below command (based on the file location on your system) on the each worker node.

For example, `chmod 600 /run/systemd/system/pf9-kubelet.service`

4.1.2 Run the below command (based on the file location on your system) on the each worker node.

For example,

`chown root:root /run/systemd/system/pf9-kubelet.service`

4.1.6 Run the below command (based on the file location on your system) on the each worker node.

For example,

`chown root:root /etc/pf9/kube.d/kubeconfigs/kubelet.yaml`

4.1.7 Run the following command to modify the file permissions of the

`--client-ca-file chmod 600 <filename>`

4.1.8 Run the following command to modify the ownership of the `--client-ca-file`.

`chown root:root <filename>`

4.1.9 Run the following command (using the config file location identified in the Audit step)

`chmod 600 /var/opt/pf9/kube/kubelet-config/bootstrap-config.yaml`

4.1.10 Run the following command (using the config file location identified in the Audit step)

`chown root:root /var/opt/pf9/kube/kubelet-config/bootstrap-config.yaml`

4.2.2 If using a Kubelet config file, edit the file to set ``authorization.mode`` to Webhook. If using executable arguments, edit the kubelet service file

`/run/systemd/system/pf9-kubelet.service` on each worker node and

set the below parameter in `KUBELET_AUTHZ_ARGS` variable.

`--authorization-mode=Webhook`

Based on your system, restart the kubelet service. For example,

`systemctl daemon-reload`

`systemctl restart kubelet.service`

4.2.7 Edit the kubelet service file `/run/systemd/system/pf9-kubelet.service`

on each worker node and remove the `--hostname-override` argument from the

`KUBELET_SYSTEM_PODS_ARGS` variable.

Based on your system, restart the kubelet service. For example,

`systemctl daemon-reload`

`systemctl restart kubelet.service`

4.2.13 Decide on an appropriate level for this parameter and set it,

either via the `--pod-max-pids` command line parameter or the `PodPidsLimit` configuration file setting.

== Summary node ==

13 checks PASS

4 checks FAIL

6 checks WARN

0 checks INFO

[INFO] 5 Kubernetes Policies

[INFO] 5.1 RBAC and Service Accounts

[WARN] 5.1.1 Ensure that the cluster-admin role is only used where required (Manual)

[WARN] 5.1.2 Minimize access to secrets (Manual)

[WARN] 5.1.3 Minimize wildcard use in Roles and ClusterRoles (Manual)

[WARN] 5.1.4 Minimize access to create pods (Manual)

[WARN] 5.1.5 Ensure that default service accounts are not actively used. (Manual)

[WARN] 5.1.6 Ensure that Service Account Tokens are only mounted where necessary (Manual)

[WARN] 5.1.7 Avoid use of `system:masters` group (Manual)

[WARN] 5.1.8 Limit use of the Bind, Impersonate and Escalate permissions in the Kubernetes cluster (Manual)

[WARN] 5.1.9 Minimize access to create persistent volumes (Manual)

[WARN] 5.1.10 Minimize access to the proxy sub-resource of nodes (Manual)

[WARN] 5.1.11 Minimize access to the approval sub-resource of certificatesigningrequests objects (Manual)

[WARN] 5.1.12 Minimize access to webhook configuration objects (Manual)

[WARN] 5.1.13 Minimize access to the service account token creation (Manual)

[INFO] 5.2 Pod Security Standards

[WARN] 5.2.1 Ensure that the cluster has at least one active policy control mechanism in place (Manual)

[WARN] 5.2.2 Minimize the admission of privileged containers (Manual)

[WARN] 5.2.3 Minimize the admission of containers wishing to share the host process ID namespace (Automated)

[WARN] 5.2.4 Minimize the admission of containers wishing to share the host IPC namespace (Automated)

[WARN] 5.2.5 Minimize the admission of containers wishing to share the host network namespace (Automated)

[WARN] 5.2.6 Minimize the admission of containers with allowPrivilegeEscalation (Automated)

[WARN] 5.2.7 Minimize the admission of root containers (Automated)

[WARN] 5.2.8 Minimize the admission of containers with the NET_RAW capability (Automated)

[WARN] 5.2.9 Minimize the admission of containers with added capabilities (Automated)

[WARN] 5.2.10 Minimize the admission of containers with capabilities assigned (Manual)

[WARN] 5.2.11 Minimize the admission of Windows HostProcess containers (Manual)

[WARN] 5.2.12 Minimize the admission of HostPath volumes (Manual)

[WARN] 5.2.13 Minimize the admission of containers which use HostPorts (Manual)

[INFO] 5.3 Network Policies and CNI

[WARN] 5.3.1 Ensure that the CNI in use supports NetworkPolicies (Manual)

[WARN] 5.3.2 Ensure that all Namespaces have NetworkPolicies defined (Manual)

[INFO] 5.4 Secrets Management

[WARN] 5.4.1 Prefer using Secrets as files over Secrets as environment variables (Manual)

[WARN] 5.4.2 Consider external secret storage (Manual)

[INFO] 5.5 Extensible Admission Control

[WARN] 5.5.1 Configure Image Provenance using ImagePolicyWebhook admission controller (Manual)

[INFO] 5.7 General Policies

[WARN] 5.7.1 Create administrative boundaries between resources using namespaces (Manual)

[WARN] 5.7.2 Ensure that the seccomp profile is set to docker/default in your Pod definitions (Manual)

[WARN] 5.7.3 Apply SecurityContext to your Pods and Containers (Manual)

[WARN] 5.7.4 The default namespace should not be used (Manual)

== Remediations policies ==

5.1.1 Identify all clusterrolebindings to the cluster-admin role. Check if they are used and if they need this role or if they could use a role with fewer privileges.

Where possible, first bind users to a lower privileged role and then remove the clusterrolebinding to the cluster-admin role :

`kubectl delete clusterrolebinding [name]`

5.1.2 Where possible, remove get, list and watch access to Secret objects in the cluster.

5.1.3 Where possible replace any use of wildcards in clusterroles and roles with specific objects or actions.

5.1.4 Where possible, remove create access to pod objects in the cluster.

5.1.5 Create explicit service accounts wherever a Kubernetes workload requires specific access to the Kubernetes API server.

Modify the configuration of each default service account to include this value

`automountServiceAccountToken: false`

5.1.6 Modify the definition of pods and service accounts which do not need to mount service account tokens to disable it.

5.1.7 Remove the `system:masters` group from all users in the cluster.

5.1.8 Where possible, remove the impersonate, bind and escalate rights from subjects.

5.1.9 Where possible, remove create access to PersistentVolume objects in the cluster.

5.1.10 Where possible, remove access to the proxy sub-resource of node objects.

5.1.11 Where possible, remove access to the approval sub-resource of certificatesigningrequest objects.

5.1.12 Where possible, remove access to the validatingwebhookconfigurations or mutatingwebhookconfigurations objects

5.1.13 Where possible, remove access to the token sub-resource of serviceaccount objects.

5.2.1 Ensure that either Pod Security Admission or an external policy control system is in place for every namespace which contains user workloads.

5.2.2 Add policies to each namespace in the cluster which has user workloads to restrict the admission of privileged containers.

5.2.3 Add policies to each namespace in the cluster which has user workloads to restrict the admission of `'hostPID'` containers.

5.2.4 Add policies to each namespace in the cluster which has user workloads to restrict the admission of `'hostIPC'` containers.

5.2.5 Add policies to each namespace in the cluster which has user workloads to restrict the admission of `'hostNetwork'` containers.

5.2.6 Add policies to each namespace in the cluster which has user workloads to restrict the admission of containers with ``.spec.allowPrivilegeEscalation`` set to `'true'`.

5.2.7 Create a policy for each namespace in the cluster, ensuring that either `'MustRunAsNonRoot'` or `'MustRunAs'` with the range of UIDs not including 0, is set.

5.2.8 Add policies to each namespace in the cluster which has user workloads to restrict the admission of containers with the `'NET_RAW'` capability.

5.2.9 Ensure that `'allowedCapabilities'` is not present in policies for the cluster unless it is set to an empty array.

5.2.10 Review the use of capabilities in applications running on your cluster. Where a namespace contains applications which do not require any Linux capabilities to operate consider adding a PSP which forbids the admission of containers which do not drop all capabilities.

5.2.11 Add policies to each namespace in the cluster which has user workloads to restrict the admission of containers that have ``.securityContext.windowsOptions.hostProcess`` set to `'true'`.

5.2.12 Add policies to each namespace in the cluster which has user workloads to restrict the admission of containers with `'hostPath'` volumes.

5.2.13 Add policies to each namespace in the cluster which has user workloads to restrict the

admission of containers which use `hostPort` sections.

5.3.1 If the CNI plugin in use does not support network policies, consideration should be given to making use of a different plugin, or finding an alternate mechanism for restricting traffic in the Kubernetes cluster.

5.3.2 Follow the documentation and create NetworkPolicy objects as you need them.

5.4.1 If possible, rewrite application code to read Secrets from mounted secret files, rather than from environment variables.

5.4.2 Refer to the Secrets management options offered by your cloud provider or a third-party secrets management solution.

5.5.1 Follow the Kubernetes documentation and setup image provenance.

5.7.1 Follow the documentation and create namespaces for objects in your deployment as you need them.

5.7.2 Use `securityContext` to enable the docker/default seccomp profile in your pod definitions.

An example is as below:

```
securityContext:
  seccompProfile:
    type: RuntimeDefault
```

5.7.3 Follow the Kubernetes documentation and apply SecurityContexts to your Pods. For a suggested list of SecurityContexts, you may refer to the CIS Security Benchmark for Docker Containers.

5.7.4 Ensure that namespaces are created to allow for appropriate segregation of Kubernetes resources and that all new resources are created in a specific namespace.

== Summary policies ==

0 checks PASS

0 checks FAIL

35 checks WARN

0 checks INFO

== Summary total ==

51 checks PASS

22 checks FAIL

58 checks WARN

0 checks INFO